



Linear Classification


Rina BUOY, PhD



ChatGPT 4.0

Disclaimer

Adopted from



6.390

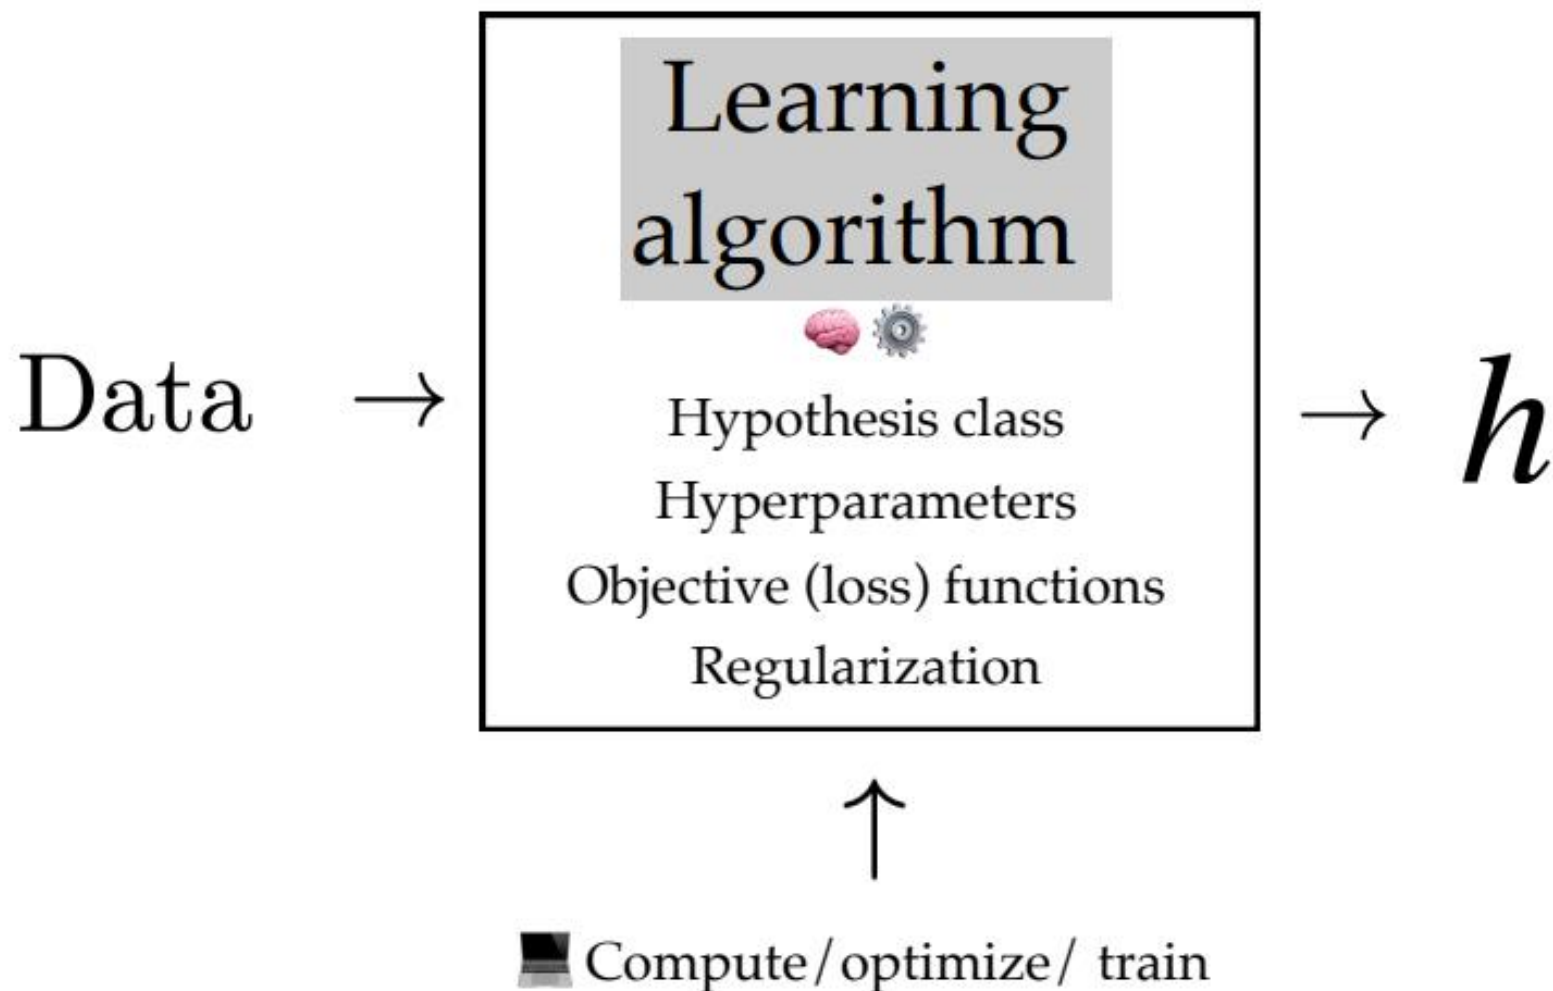
**Introduction to Machine Learning
(Fall 2024)**

<https://introml.mit.edu/fall24>

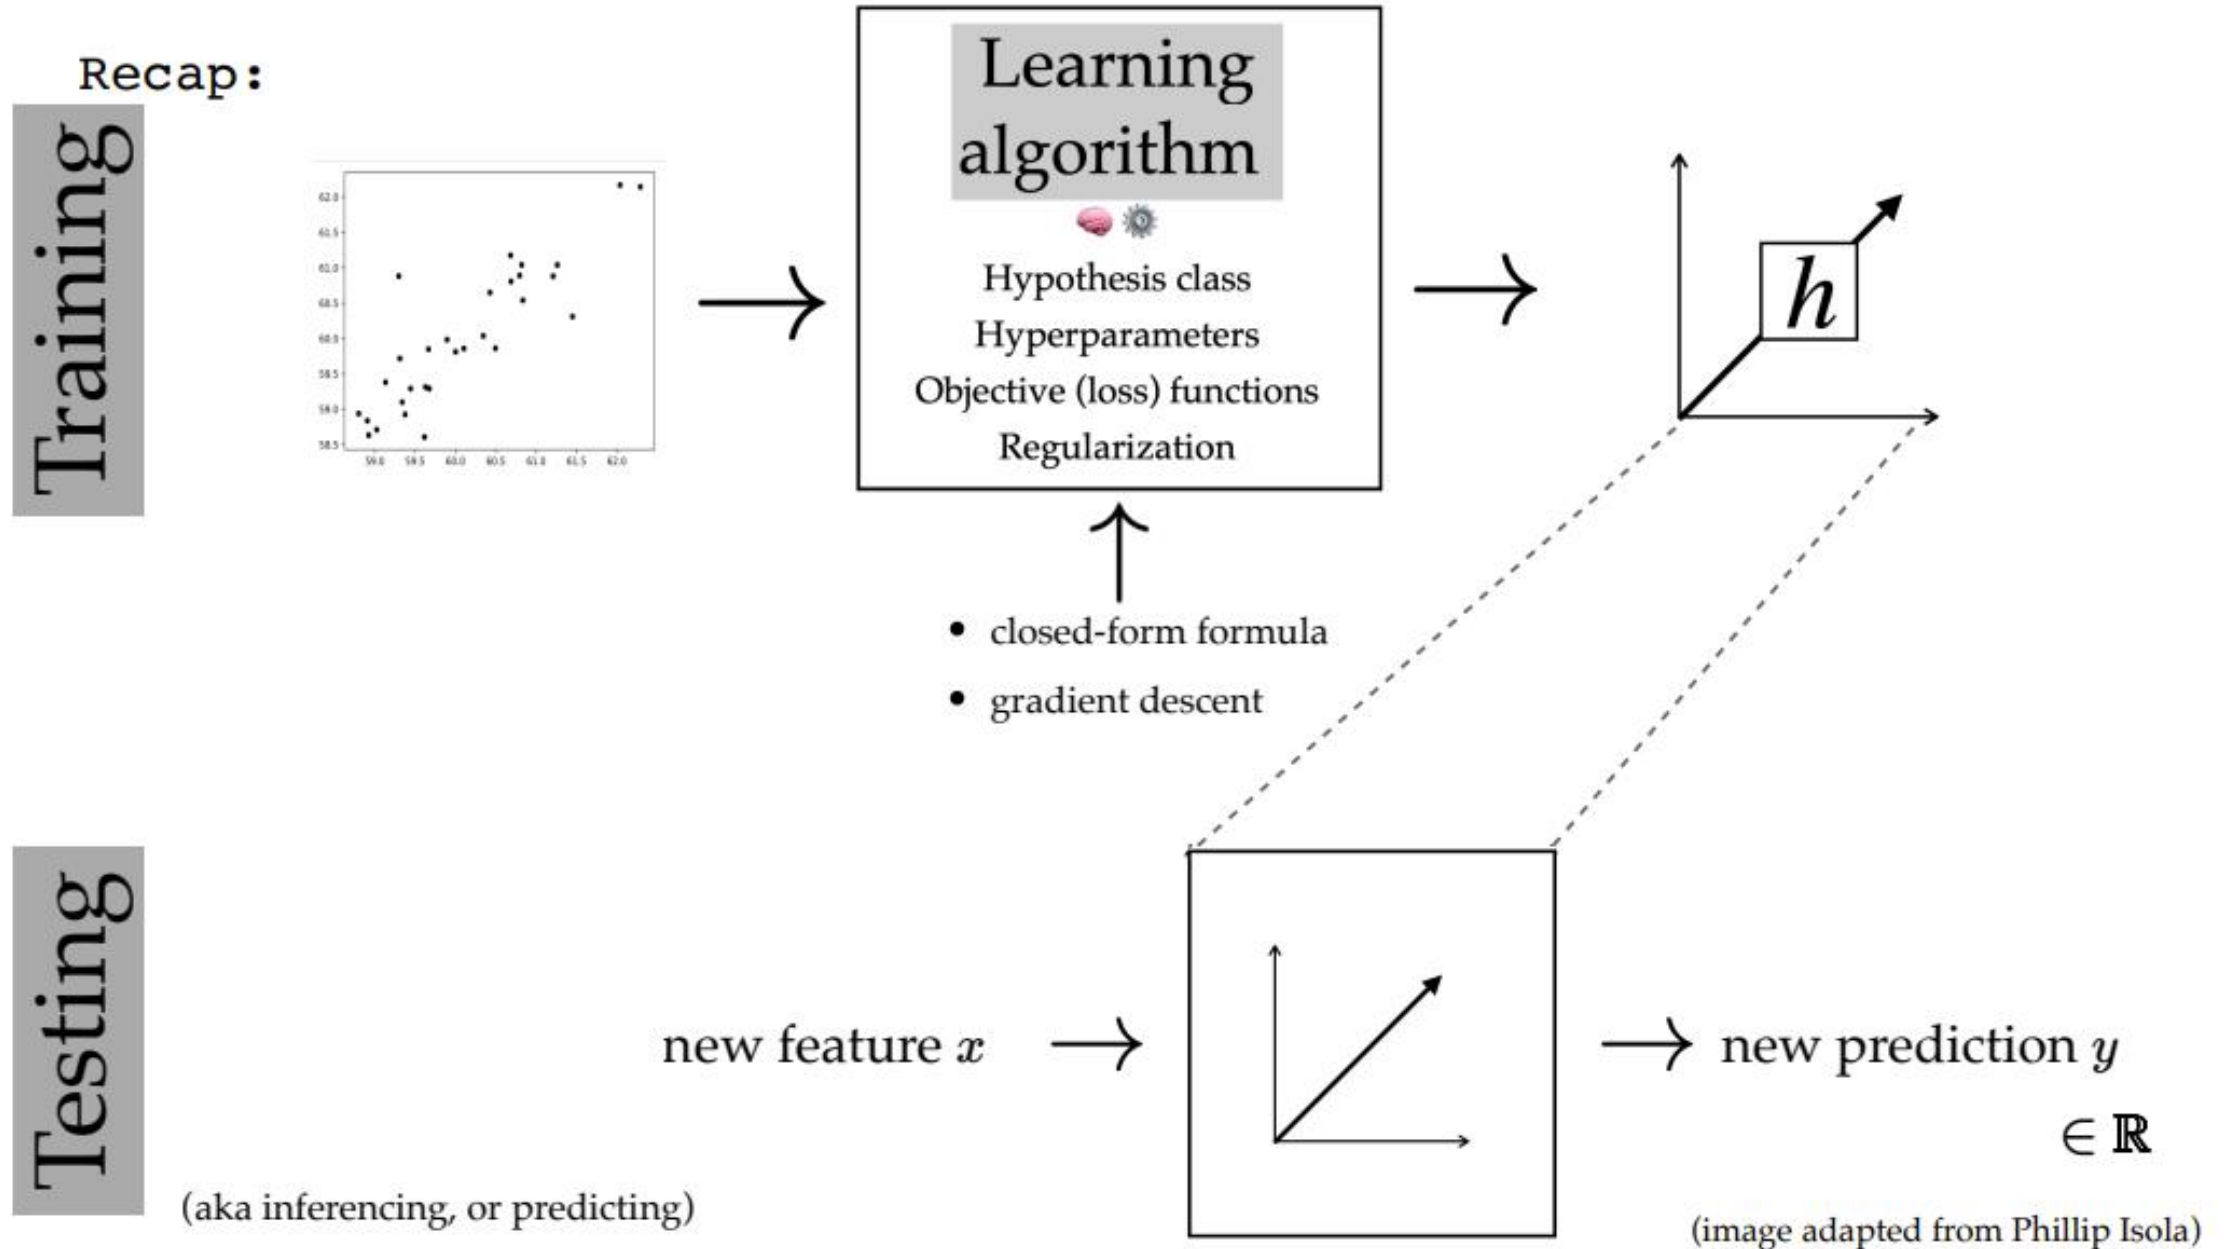
Outline

- **Recap, classification setup**
- Linear classifiers
 - Separator, normal vector, and separability
- Linear logistic classifiers
 - Motivation, sigmoid, and negative log-likelihood loss
- Multi-class classifiers
 - One-hot encoding, softmax, and cross-entropy

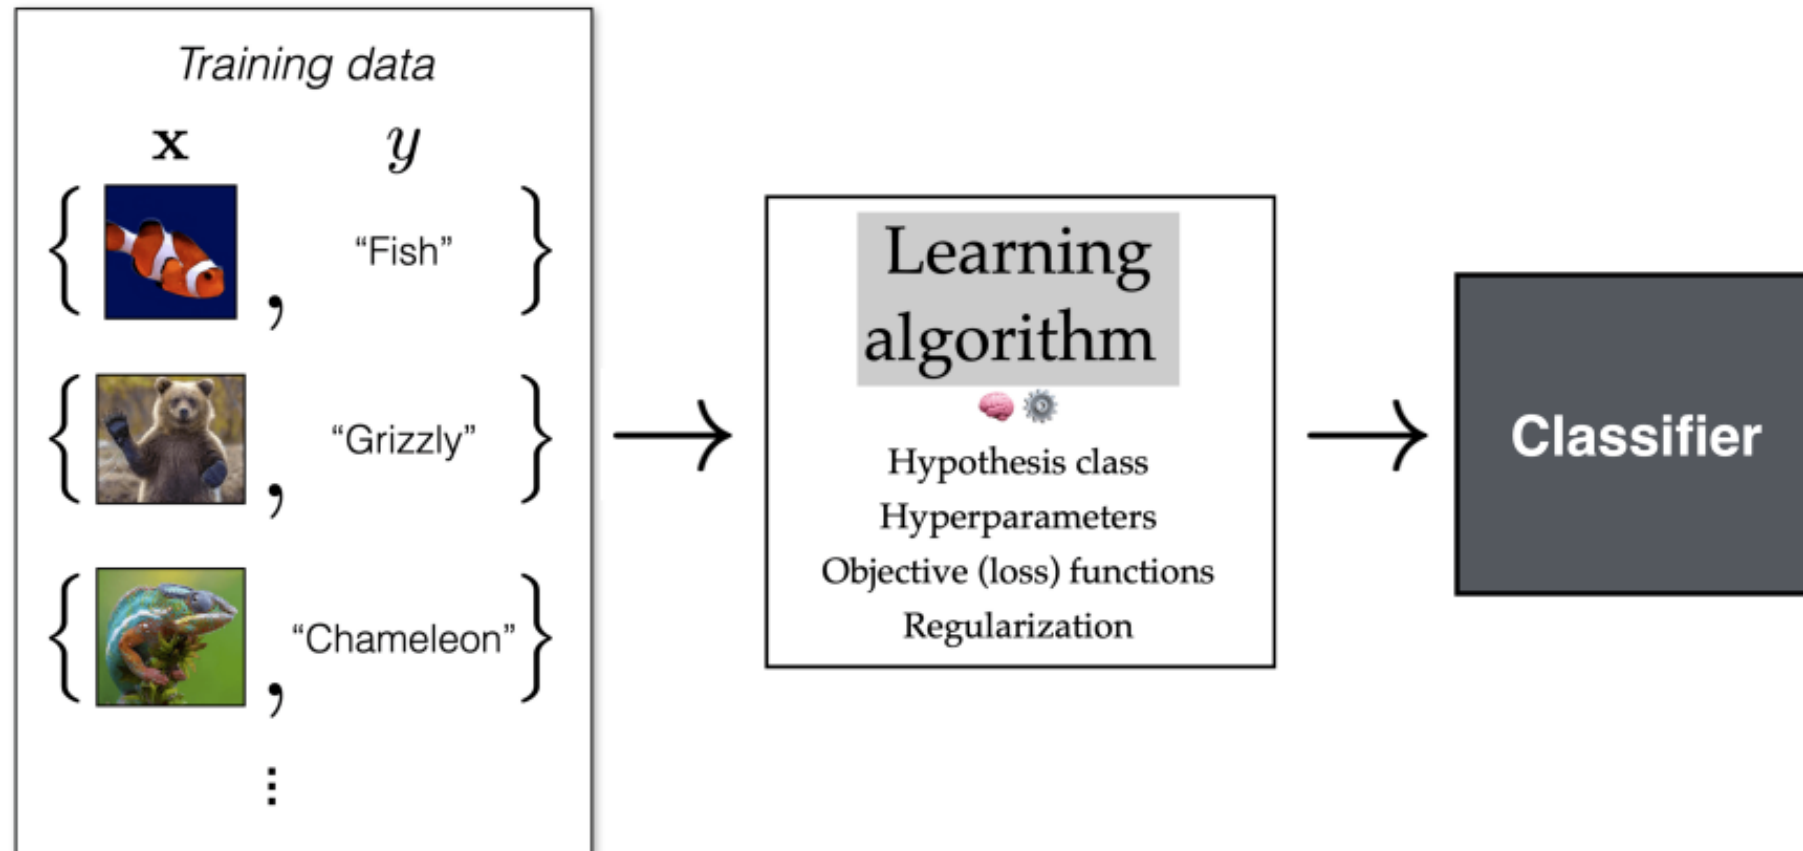
Recap:



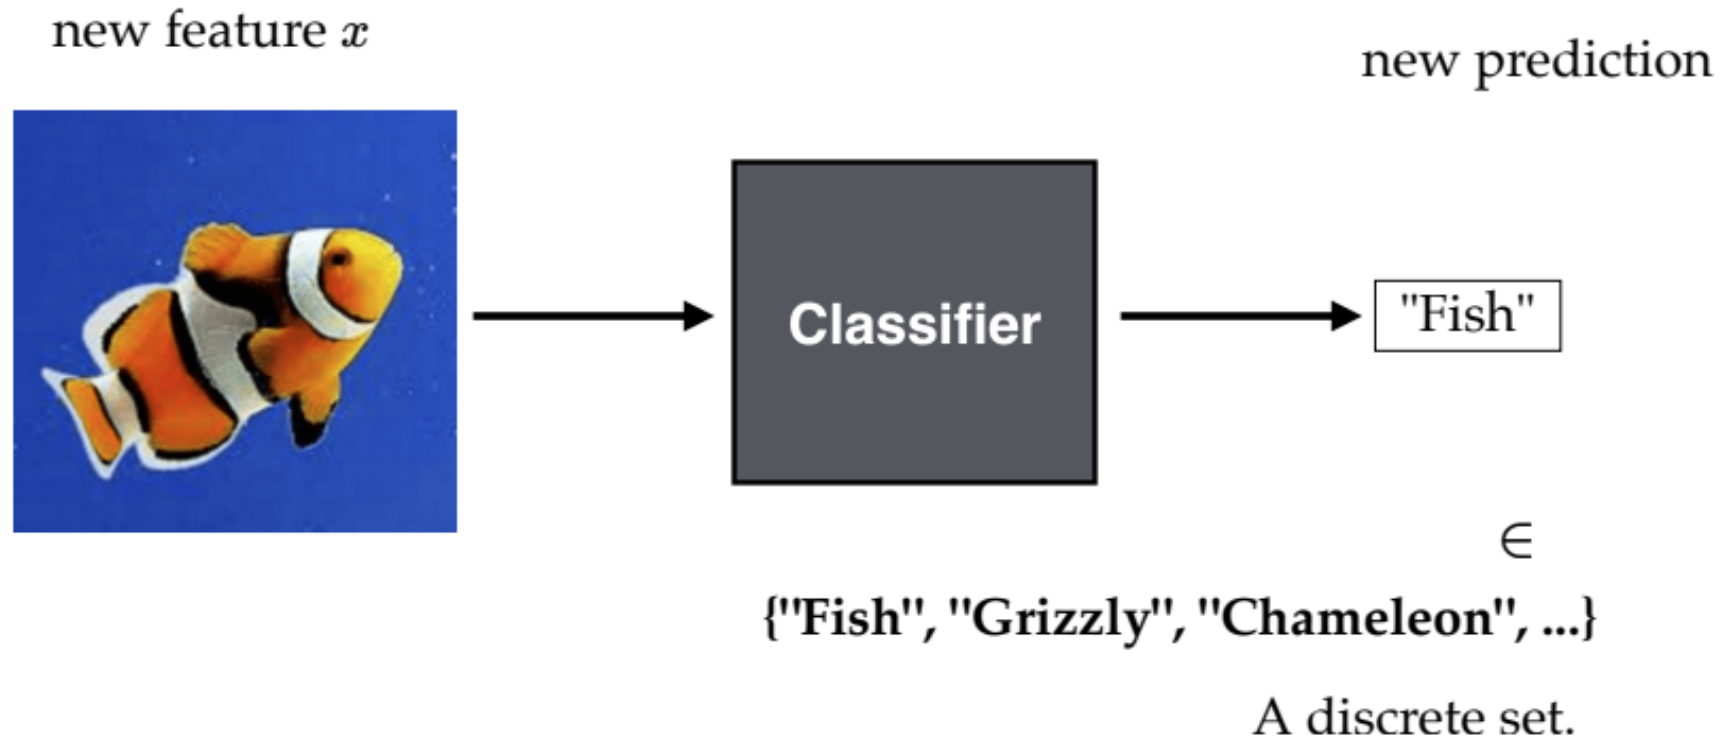
(image adapted from Phillip Isola)



Classification Setup




(image adapted from Phillip Isola)



(image adapted from Phillip Isola)

Linear Regression from Probability Perspective



Maximum Log-Likelihood

$$\hat{\theta} = \operatorname{argmax}_{\theta} \log p(\mathcal{D} \mid \theta)$$

Maximum Log-Likelihood

$$\hat{\theta} = \operatorname{argmax}_{\theta} \log p(\mathcal{D} \mid \theta)$$

$$\begin{aligned} l(\theta) &:= \log p(\mathcal{D} \mid \theta) \\ &= \log \left(\prod_{i=1}^N p(y_i \mid \mathbf{x}_i, \theta) \right) \\ &= \sum_{i=1}^N \log p(y_i \mid \mathbf{x}_i, \theta) \end{aligned}$$

Negative Log-Likelihood

$$\text{NLL}(\theta) = - \sum_{i=1}^N \log p(y_i \mid \mathbf{x}_i, \theta)$$

Negative Log-Likelihood

$$\begin{aligned}\text{NLL}(\theta) &= - \sum_{i=1}^N \log p(y_i \mid \mathbf{x}_i, \theta) \\ &= - \sum_{i=1}^N \log \left[\left(\frac{1}{2\pi\sigma^2} \right)^{\frac{1}{2}} \exp \left(-\frac{1}{2\sigma^2} (y_i - \beta^T \mathbf{x}_i)^2 \right) \right] \\ &= - \sum_{i=1}^N \frac{1}{2} \log \left(\frac{1}{2\pi\sigma^2} \right) - \frac{1}{2\sigma^2} (y_i - \beta^T \mathbf{x}_i)^2 \\ &= - \frac{N}{2} \log \left(\frac{1}{2\pi\sigma^2} \right) - \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \beta^T \mathbf{x}_i)^2 \\ &= - \frac{N}{2} \log \left(\frac{1}{2\pi\sigma^2} \right) - \frac{1}{2\sigma^2} \text{RSS}(\beta)\end{aligned}$$

Optimal Parameters

$$\frac{\partial NLL}{\partial \beta} = 0 \quad \Downarrow$$

$$\hat{\beta}_{\text{OLS}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Maximum Likelihood Formulation

1. Choose a suitable probability distribution $Pr(\mathbf{y}|\boldsymbol{\theta})$ that is defined over the domain of the predictions \mathbf{y} and has distribution parameters $\boldsymbol{\theta}$.
2. Set the machine learning model $\mathbf{f}[\mathbf{x}, \phi]$ to predict one or more of these parameters so $\boldsymbol{\theta} = \mathbf{f}[\mathbf{x}, \phi]$ and $Pr(\mathbf{y}|\boldsymbol{\theta}) = Pr(\mathbf{y}|\mathbf{f}[\mathbf{x}, \phi])$.
3. To train the model, find the network parameters $\hat{\phi}$ that minimize the negative log-likelihood loss function over the training dataset pairs $\{\mathbf{x}_i, \mathbf{y}_i\}$:

$$\hat{\phi} = \underset{\phi}{\operatorname{argmin}} [L[\phi]] = \underset{\phi}{\operatorname{argmin}} \left[- \sum_{i=1}^I \log \left[Pr(\mathbf{y}_i | \mathbf{f}[\mathbf{x}_i, \phi]) \right] \right]. \quad (5.7)$$

Outline


- Recap, classification setup
- Linear classifiers
 - Separator, normal vector, and separability
- Linear logistic classifiers
 - Motivation, sigmoid, and negative log-likelihood loss
- Multi-class classifiers
 - One-hot encoding, softmax, and cross-entropy

(vanilla, sign-based, binary) Linear Classifier

- Each data point:
 - features $[x_1, x_2, \dots, x_d]$
 - label $y \in \{\text{positive, negative}\}$ (or $\{\text{dog, cat}\}$, $\{\text{pizza, not pizza}\}$, $\{+1, 0\}$)
- A (vanilla, sign-based, binary) linear classifier is parameterized by $[\theta_1, \theta_2, \dots, \theta_d, \theta_0]$
- To *use* a given classifier make prediction:
 - do linear combination: $z = (\theta_1 x_1 + \theta_2 x_2 + \dots + \theta_d x_d) + \theta_0$
 - predict positive label if $z > 0$, otherwise, negative label.

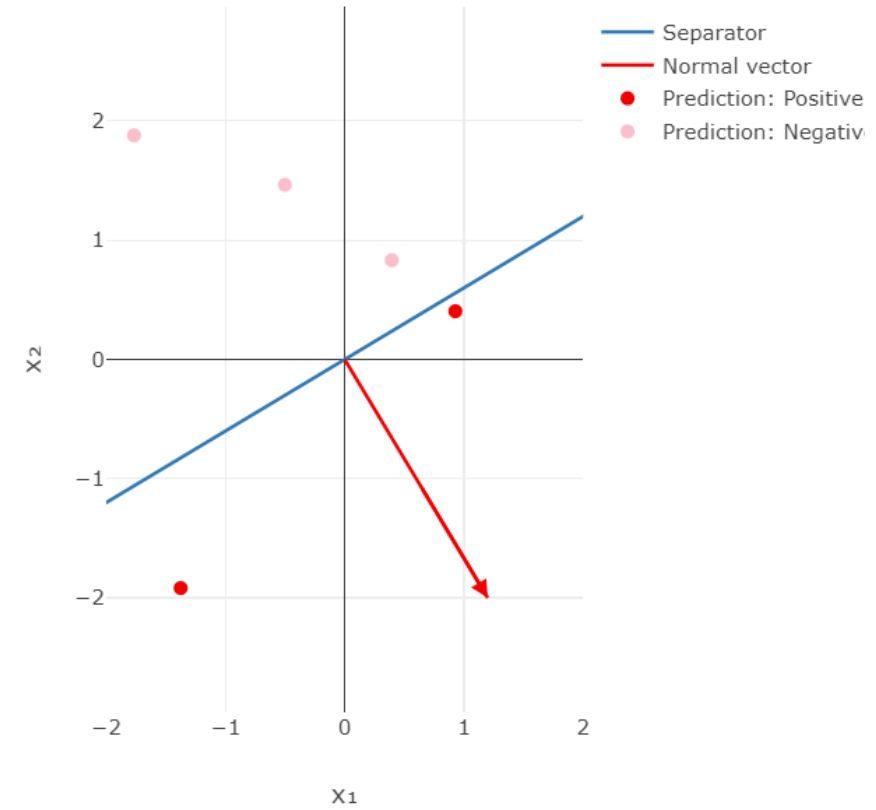
View of the feature space (x_1 and x_2) and decision helper (z)

$$z = \theta_1 x_1 + \theta_2 x_2 + \theta_0$$



- Prediction: Positive
- Prediction: Negative

View of the feature space (x_1 and x_2)



θ_1 :



1.2

θ_2 :



-2.0

θ_0 :



0.0


Toggle $z=0$ Surface

- Now let's try to *learn* a linear classifier
- One natural loss: $\mathcal{L}_{01}(g, a) = \begin{cases} 0 & \text{if guess} = \text{actual} \\ 1 & \text{otherwise} \end{cases}$
- Combined with the linear classifier hypothesis:


$$\mathcal{L}_{01}(x^{(i)}, y^{(i)}; \theta, \theta_0) = \begin{cases} 0 & \text{if } \text{sign}(\theta^\top x^{(i)} + \theta_0) = y^{(i)} \\ 1 & \text{otherwise} \end{cases}$$

- Very intuitive, and easy to evaluate 🥰
 - Induced concept: separability
- Very hard to optimize (NP-hard) 😞
 - "Flat" almost everywhere (zero gradient)
 - "Jumps" elsewhere (no gradient)

Demo dataset



Sum of 0-1 loss (on the demo dataset on the left)



Try to draw the separator and normal vector given by $(\theta_1 = -1$, and $\theta_2 = 1$) on the 2D plot, and make sense of the loss given in the 3D plot.

Outline

- Recap, classification setup
- Linear classifiers
 - Separator, normal vector, and separability
- Linear logistic classifiers
 - Motivation, sigmoid, and negative log-likelihood loss
- Multi-class classifiers
 - One-hot encoding, softmax, and cross-entropy

Linear Logistic Classifier


- Mainly motivated to address the gradient issue in *learning* a "vanilla" linear classifier
 - The gradient issue is caused by both the 0/1 loss, and the sign functions nested in.

$$\mathcal{L}_{01}(x^{(i)}, y^{(i)}; \theta, \theta_0) = \begin{cases} 0 & \text{if } \text{sign}(\theta^\top x^{(i)} + \theta_0) = y^{(i)} \\ 1 & \text{otherwise} \end{cases}$$

- But has nice probabilistic interpretation too.
- As before, let's first look at how to make prediction with a *given* linear logistic classifier

(Binary) Linear Logistic Classifier

- Each data point:
 - features $[x_1, x_2, \dots, x_d]$
 - label $y \in \{\text{positive, negative}\}$
- A (binary) linear **logistic** classifier is parameterized by $[\theta_1, \theta_2, \dots, \theta_d, \theta_0]$
- To *use* a given classifier make prediction:
 - do linear combination: $z = (\theta_1 x_1 + \theta_2 x_2 + \dots + \theta_d x_d) + \theta_0$
 - predict positive label if



$$\sigma(z) = \sigma(\theta^\top x + \theta_0) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + e^{-(\theta^\top x + \theta_0)}} > 0.5$$

otherwise, negative label.


Sigmoid : a smooth step function

$$\sigma(z) = \sigma(\theta^\top x + \theta_0) = \frac{1}{1 + e^{-(\theta^\top x + \theta_0)}}$$

- "sandwiched" between 0 and 1
vertically (never 0 or 1 mathematically)
- θ, θ_0 can flip, squeeze, expand, shift *horizontally*
- $\sigma(\cdot)$ interpreted as the *probability / confidence* that feature x has positive label. Predict positive if

$$\sigma(z) = \sigma(\theta^\top x + \theta_0) > 0.5$$

- monotonic, very nice / elegant gradient (see recitation/hw)



$$\sigma(z) = \sigma(\theta^\top x + \theta_0) = \frac{1}{1 + e^{-(\theta^\top x + \theta_0)}} \quad \Rightarrow \quad \textbf{Probability}$$

$$\theta^\top x + \theta_0 \quad \Rightarrow \quad \textbf{Logit or log odd} \quad \Rightarrow \quad \log\left(\frac{p}{1-p}\right)$$



$$\frac{p}{1-p} \quad \Rightarrow \quad \textbf{Odd ratio}$$

If the probability of rain is 0.5, the odd ratio of rain to no rain is 1 and the log odd is 0.



e.g. suppose, wanna predict whether to bike to school.

with **given** parameters, how do I make prediction?

1 feature:
$$g(x) = \sigma(\theta x + \theta_0)$$
$$= \frac{1}{1 + \exp\{- (\theta x + \theta_0)\}}$$

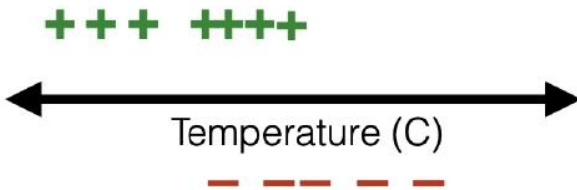


2 features:
$$g(x) = \sigma(\theta^\top x + \theta_0)$$
$$= \frac{1}{1 + \exp\{- (\theta^\top x + \theta_0)\}}$$





(image credit: Tamara Broderick)

Learning a logistic regression classifier

training data: 

$$g(x) = \sigma(\theta x + \theta_0)$$




- Let the labels $y \in \{+1, 0\}$

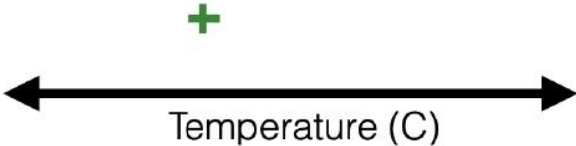
$$\mathcal{L}_{\text{nl}}(\text{guess}, \text{actual})$$

$$= -[\text{actual} \cdot \log(\text{guess}) + (1 - \text{actual}) \cdot \log(1 - \text{guess})]$$

$$= -\left[y^{(i)} \log g^{(i)} + (1 - y^{(i)}) \log (1 - g^{(i)})\right]$$






If $y^{(i)} = 1$

training data: 

$\mathcal{L}_{\text{nll}}(\text{guess}, \text{actual})$

$$= - \left[y^{(i)} \log g^{(i)} + (1 - y^{(i)}) \log (1 - g^{(i)}) \right]$$

$$g(x) = \sigma(\theta x + \theta_0)$$



If $y^{(i)} = 0$




training
data:



$\mathcal{L}_{\text{nll}}(\text{guess}, \text{actual})$

$$= - \left[\cancel{y^{(i)} \log g^{(i)}} + \cancel{(1 - y^{(i)}) \log (1 - g^{(i)})} \right]$$

$$g(x) = \sigma(\theta x + \theta_0)$$




Logistic Regression


- Minimize using negative-log-likelihood loss:


$$J_{lr} = \frac{1}{n} \sum_{i=1}^n \mathcal{L}_{\text{nll}} \left(\sigma \left(\theta^\top x^{(i)} + \theta_0 \right), y^{(i)} \right)$$

- Convex, differentiable with **nice** (elegant) gradients
- Doesn't have a closed-form solution
- Can still run gradient descent
- But, a gotcha: when training data is linearly separable



$$g(x) = \sigma \left(\theta^T x + \theta_0 \right)$$








http://www.statistics4u.com/fundstat_eng/cc_data_structur
e.html

Regularized Logistic Regression

$$J(\theta, \theta_0) = \left(\frac{1}{n} \sum_{i=1}^n \mathcal{L}_{\text{nll}} \left(\sigma \left(\theta^\top x^{(i)} + \theta_0 \right), y^{(i)} \right) \right) + \lambda \|\theta\|^2$$

- $\lambda \geq 0$
- No regularizing θ_0 (think: why?)
- Penalizes being overly certain
- Objective is still differentiable and convex (gradient descent)



L1 Regularization

$$\text{Cost} = \sum_{i=0}^N (y_i - \sum_{j=0}^M x_{ij} W_j)^2 + \lambda \sum_{j=0}^M |W_j|$$

L2 Regularization


$$\text{Cost} = \underbrace{\sum_{i=0}^N (y_i - \sum_{j=0}^M x_{ij} W_j)^2}_{\text{Loss function}} + \lambda \underbrace{\sum_{j=0}^M W_j^2}_{\text{Regularization Term}}$$

Outline

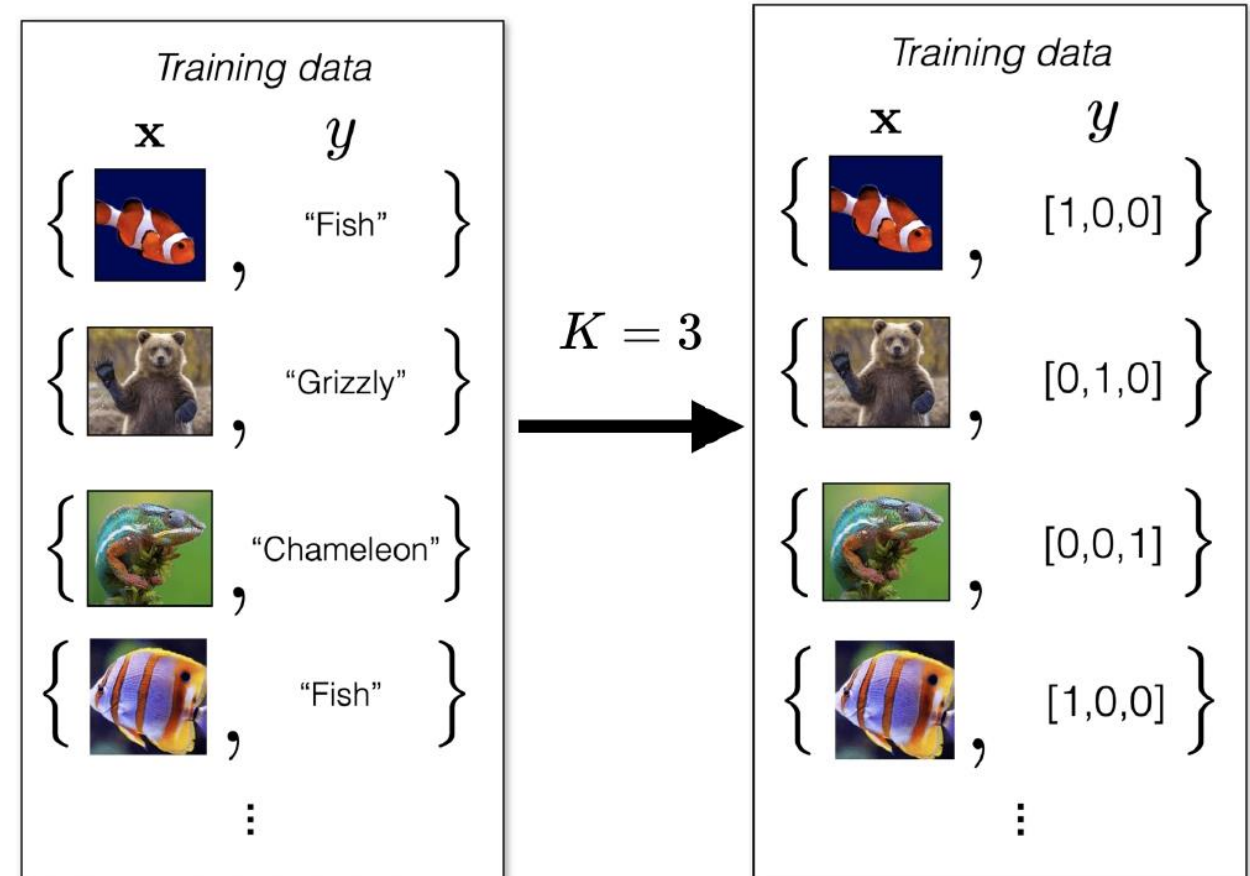
- Recap, classification setup
- Linear classifiers
 - Separator, normal vector, and separability
- Linear logistic classifiers
 - Motivation, sigmoid, and negative log-likelihood loss
- Multi-class classifiers
 - One-hot encoding, softmax, and cross-entropy

One-hot labels

- Generalizes from binary labels
- Suppose K classes



One-hot encoding



(image adapted from Phillip Isola)

Softmax

- Generalizes sigmoid
- Applied on z element-wise

Two classes

$$\uparrow$$
$$z = \theta^\top x + \theta_0$$

scalar

$$\uparrow$$
$$g = \sigma(z) = \frac{1}{1 + \exp(-z)}$$

scalar

K classes

$$\uparrow$$
$$z = \theta^\top x + \theta_0$$

K -by-1 Vector

$$\uparrow$$
$$g = \text{softmax}(z) = \begin{bmatrix} \exp(z_1) / \sum_i \exp(z_i) \\ \vdots \\ \exp(z_K) / \sum_i \exp(z_i) \end{bmatrix}$$

K -by-1 Vector

Negative log-likelihood multi-class loss

- Generalizes negative log likelihood loss
- Also known as cross-entropy

Two classes


$$\mathcal{L}_{\text{nll}}(\mathbf{g}, \mathbf{y}) = -(y \log g + (1 - y) \log (1 - g))$$


- Appears as sum of two terms
- Only one term "activates" for a single data point

K classes

$$\mathcal{L}_{\text{nllm}}(\mathbf{g}, \mathbf{y}) = - \sum_{k=1}^K y_k \cdot \log (g_k)$$

- Appears as sum of K terms
- Only one term "activates" for a single data point






(image adapted from Phillip Isola)

Classification


Image classification played a pivotal role in kicking off the current wave of AI enthusiasm

ImageNet Classification 2012

- Krizhevsky et al. -- 16.4% error (top-5)
- Next best (non-convnet) – 26.2% error



What comes next?



Summary

- Classification: a supervised learning problem, similar to regression, but where the output/label is in a discrete set.
- Binary classification: only two possible label values.
- Linear binary classification: think of θ and θ_0 as defining a $d-1$ dimensional hyperplane that **cuts** the d -dimensional feature space into two half-spaces.
- 0-1 loss: a natural loss function for classification, BUT, hard to optimize.
- Sigmoid function: motivation and properties.
- Negative-log-likelihood loss: smoother and has nice probabilistic motivations. We can optimize via (S)GD.
- Regularization is still important.
- The generalization to multi-class via (one-hot encoding, and softmax mechanism)
- Other ways to generalize to multi-class (see hw/lab)