



Practical Application in Machine Learning

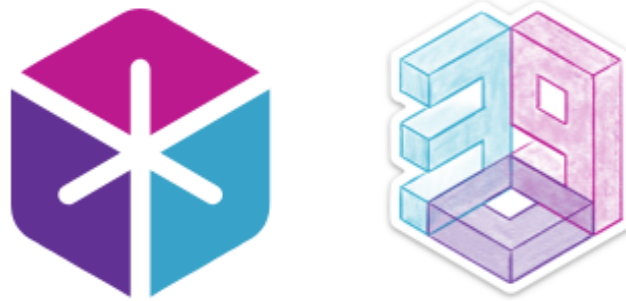
Rina BUOY, PhD



ChatGPT 4.0

Disclaimer

Adopted from



6.390

**Introduction to Machine Learning
(Fall 2024)**

<https://introml.mit.edu/fall24>

Expected prerequisite background

Things we expect you to know (we use these constantly, but don't teach them explicitly):

Programming (e.g. as in 6.101[009] or 6.121[006])

- Intermediate Python, including classes
- Exposure to algorithms – ability to understand & discuss pseudo-code, and implement in Python

Linear Algebra (e.g. as in 18.06, 18.C06, 18.03, or 18.700)

- Matrix manipulations: transpose, multiplication, inverse etc.
- Points and planes in high-dimensional space
- (Together with calculus): taking gradients, matrix calculus

Useful background

Things it helps to have prior exposure to, but we don't expect (we use these in 6.390, but will discuss as we go):

- numpy (Python package for matrix/linear algebra)
- pytorch (python package for modern ml models like deep neural networks)
- Basic discrete probability: random variables, independence, conditioning

What we're teaching: Machine Learning!

Given:

- a **collection of examples** (gene sequences, documents, ...)
- an **encoding of those examples** in a computer (as vectors)



Derive:

- a **computational model** (called a hypothesis) that describes relationships within and among the examples that is expected to characterize well new examples from that same population, to make good predictions or decisions

A model might:

- **classify images** of cells as to whether they're cancerous
- **specify groupings (clusters)** of documents that address similar topics
- **steer** a car appropriately given lidar images of the surroundings

Very roughly, ML can be categorized into



(the categorization can be refined, e.g. there are active learning, semi-supervised, selective, contrastive, few-shot, inverse reinforcement learning...)

[Slides adapted from 6.790]

Supervised learning

Goal: correctly classify so far unseen test images



Goal: predict to what degree a drug candidate binds to the intended target protein (based on a dataset of already-screened molecules against the target)



- Learning a machine translation system from pairs of sentences

Spanish (input)

Aquí tienes un bolígrafo

Las conferencias de ML son divertidas

Todo el mundo debería estudiar AI

...

English (output)

Here's a pen

ML conferences are fun

Everyone should study AI

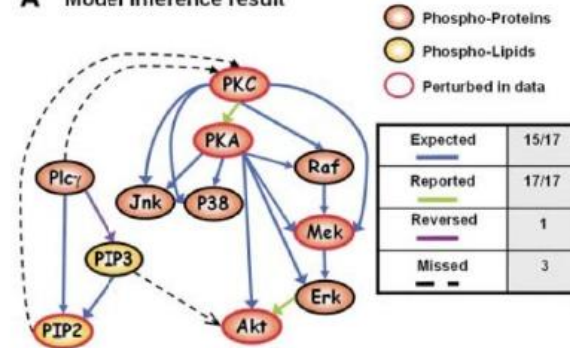
...

[Slides adapted from 6.790]

Unsupervised learning



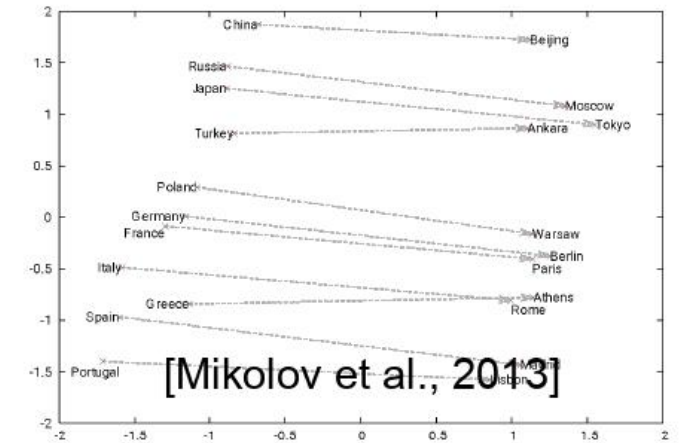
dependency
/causal
structure

A Model inference result



[Sachs et al 05]

dimensionality reduction, embedding




[courtesy of Jason Yim]

Over 3D protein structures, etc.

**+Self-Supervised
paradigm**

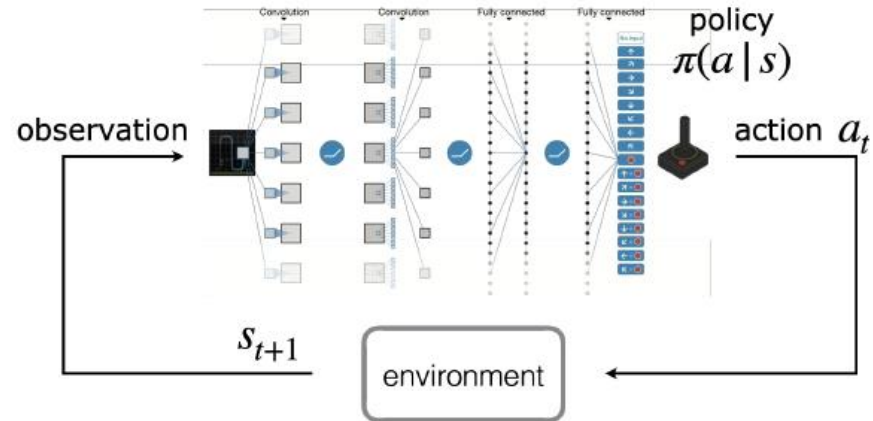
de-noising diffusion models over images



[image from
Rissanen et al 2022]

[Slides adapted from 6.790]

Reinforcement learning



Step 1
Collect demonstration data and train a supervised policy.

A prompt is sampled from our prompt dataset.

A labeler demonstrates the desired output behavior.

This data is used to fine-tune GPT-3.5 with supervised learning.




Step 2
Collect comparison data and train a reward model.

A prompt and several model outputs are sampled.

A labeler ranks the outputs from best to worst.

This data is used to train our reward model.



Step 3
Optimize a policy against the reward model using the PPO reinforcement learning algorithm.



A new prompt is sampled from the dataset.

The PPO model is initialized from the supervised policy.

The policy generates an output.

The reward model calculates a reward for the output.

The reward is used to update the policy using PPO.



Single-player game played by AlphaTensor, where the goal is to find a correct matrix multiplication algorithm. The state of the game is a cubic array of numbers (shown as grey for 0, blue for 1, and green for -1), representing the remaining work to be done.

[Slides adapted from 6.790]

Machine learning (ML): why & what

- **What is ML?**
 - Roughly, a set of methods for making predictions and decisions from data.
- **Why study ML?**
 - To apply; to understand; to evaluate; to create
- **What do we have?**
 - Data! And computation!
- **What do we want?**
 - To make predictions on new data!
- **How do we learn to make those decisions?**
 - The topic of this course!

What do we have?

- There are many different **problem classes** in ML
 - We will first focus on an instance of **supervised learning** known as **regression**.

(Training) data


- n **training data** points
- For data point $i \in \{1, \dots, n\}$

- **Feature vector**

$$x^{(i)} = (x_1^{(i)}, \dots, x_d^{(i)})^\top \in \mathbb{R}^d$$

- **Label** $y^{(i)} \in \mathbb{R}$


- **Training data** $\mathcal{D}_n = \{(x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)})\}$




What do we want?

We want a “good” way to label new feature vectors

- How to label? Learn a hypothesis
- We typically consider a class of possible hypotheses



Input:
Feature vector




Output:
Label

how well our hypothesis labels new feature vectors depends largely on how expressive the hypothesis class is

Warning: Overfitting vs. Generalization

What we really want is to generalize to **future data**!

- What we don't want:
 - Model does not capture the input-output relationship
→ **Underfitting**
 - Model too specialized to training data → **Overfitting**




What do we want?

We may consider the class of **linear regressors**:

- Hypotheses take the form:

$$h(x; \underbrace{\theta, \theta_0}) = \theta^\top x + \theta_0$$

Generally, we might refer to the set of all learned parameters as Θ (capital θ)



How **good** is a hypothesis?

Hopefully predict well on **future data**


How good is a regressor at one point?

- Quantify the error using a **loss function**, $\mathcal{L}(g, a)$

g : guess
 a : actual

- Common choice: squared loss:

$$\mathcal{L}(g, a) = (g - a)^2$$



h : hypothesis function (outputs g)
 x : input, θ : parameters, y : actual

- Training error:** $\mathcal{E}_n(h; \Theta) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(h(x^{(i)}; \Theta), y^{(i)})$


- Validation or Test error** (n' new points):

$$\mathcal{E}(h) = \frac{1}{n'} \sum_{i=n+1}^{n+n'} \mathcal{L}(h(x^{(i)}), y^{(i)})$$

How do we learn?

- Have data; have hypothesis class
- Want to choose (learn) a good hypothesis (a set of parameters)

What we want:



How to get it:
(Next time!)

